

Activity__ Course 5 Automatidata project lab

March 13, 2024

1 Automatidata project

Course 5 - Regression Analysis: Simplify complex data relationships

The data consulting firm Automatidata has recently hired you as the newest member of their data analytics team. Their newest client, the NYC Taxi and Limousine Commission (New York City TLC), wants the Automatidata team to build a multiple linear regression model to predict taxi fares using existing data that was collected over the course of a year. The team is getting closer to completing the project, having completed an initial plan of action, initial Python coding work, EDA, and A/B testing.

The Automatidata team has reviewed the results of the A/B testing. Now it's time to work on predicting the taxi fare amounts. You've impressed your Automatidata colleagues with your hard work and attention to detail. The data team believes that you are ready to build the regression model and update the client New York City TLC about your progress.

A notebook was structured and prepared to help you in this project. Please complete the following questions.

2 Course 5 End-of-course project: Build a multiple linear regression model

In this activity, you will build a multiple linear regression model. As you've learned, multiple linear regression helps you estimate the linear relationship between one continuous dependent variable and two or more independent variables. For data science professionals, this is a useful skill because it allows you to consider more than one variable against the variable you're measuring against. This opens the door for much more thorough and flexible analysis to be completed.

Completing this activity will help you practice planning out and building a multiple linear regression model based on a specific business need. The structure of this activity is designed to emulate the proposals you will likely be assigned in your career as a data professional. Completing this activity will help prepare you for those career moments.

The purpose of this project is to demonstrate knowledge of EDA and a multiple linear regression model

The goal is to build a multiple linear regression model and evaluate the model *This activity has three parts:*

Part 1: EDA & Checking Model Assumptions * What are some purposes of EDA before constructing a multiple linear regression model?

Part 2: Model Building and evaluation * What resources do you find yourself using as you complete this stage?

Part 3: Interpreting Model Results

- What key insights emerged from your model(s)?
- What business recommendations do you propose based on the models built?

3 Build a multiple linear regression model

4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

4.1.1 Task 1. Imports and loading

Import the packages that you've learned are needed for building linear regression models.

```
[1]: # Imports
# Packages for numerics + dataframes
### YOUR CODE HERE ###
import pandas as pd
import numpy as np

# Packages for visualization
### YOUR CODE HERE ###
import matplotlib.pyplot as plt
import seaborn as sns

# Packages for date conversions for calculating trip durations
### YOUR CODE HERE ###
from datetime import datetime
from datetime import date
from datetime import timedelta

# Packages for OLS, MLR, confusion matrix
```

```

### YOUR CODE HERE ###
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import sklearn.metrics as metrics # For confusion matrix
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error

```

Note: Pandas is used to load the NYC TLC dataset. As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```

[2]: # Load dataset into dataframe
df0=pd.read_csv("2017_Yellow_Taxi_Trip_Data.csv")

```

4.2 PACE: Analyze

In this stage, consider the following question where applicable to complete your code response:

- What are some purposes of EDA before constructing a multiple linear regression model?

==> ENTER YOUR RESPONSE HERE

4.2.1 Task 2a. Explore data with EDA

Analyze and discover data, looking for correlations, missing data, outliers, and duplicates.

Start with `.shape` and `.info()`.

```

[3]: # Start with `.shape` and `.info()`
### YOUR CODE HERE ###
df = df0.copy()
print(df.shape)
df.info()

```

```

(22699, 18)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            22699 non-null  int64
1   VendorID              22699 non-null  int64
2   tpep_pickup_datetime  22699 non-null  object
3   tpep_dropoff_datetime 22699 non-null  object
4   passenger_count       22699 non-null  int64
5   trip_distance         22699 non-null  float64
6   RatecodeID           22699 non-null  int64

```

```

7   store_and_fwd_flag      22699 non-null object
8   PULocationID            22699 non-null int64
9   DOLocationID            22699 non-null int64
10  payment_type             22699 non-null int64
11  fare_amount              22699 non-null float64
12  extra                    22699 non-null float64
13  mta_tax                  22699 non-null float64
14  tip_amount               22699 non-null float64
15  tolls_amount             22699 non-null float64
16  improvement_surcharge    22699 non-null float64
17  total_amount             22699 non-null float64

```

dtypes: float64(8), int64(7), object(3)

memory usage: 3.1+ MB

Check for missing data and duplicates using `.isna()` and `.drop_duplicates()`.

```

[4]: # Check for missing data and duplicates using .isna() and .drop_duplicates()
    ### YOUR CODE HERE ###
    print("Missing values: ", df.isna().any(axis = 1).sum())
    print('Shape of dataframe with duplicates dropped:', df.drop_duplicates().shape)

```

Missing values: 0

Shape of dataframe with duplicates dropped: (22699, 18)

Use `.describe()`.

```

[5]: # Use .describe()
    ### YOUR CODE HERE ###
    df.describe()

```

```

[5]:      Unnamed: 0      VendorID  passenger_count  trip_distance  \
count  2.269900e+04  22699.000000      22699.000000      22699.000000
mean    5.675849e+07      1.556236          1.642319          2.913313
std     3.274493e+07      0.496838          1.285231          3.653171
min     1.212700e+04      1.000000          0.000000          0.000000
25%     2.852056e+07      1.000000          1.000000          0.990000
50%     5.673150e+07      2.000000          1.000000          1.610000
75%     8.537452e+07      2.000000          2.000000          3.060000
max     1.134863e+08      2.000000          6.000000          33.960000

      RatecodeID  PULocationID  DOLocationID  payment_type  fare_amount  \
count  22699.000000  22699.000000  22699.000000  22699.000000  22699.000000
mean     1.043394    162.412353    161.527997     1.336887     13.026629
std     0.708391     66.633373     70.139691     0.496211     13.243791
min     1.000000     1.000000     1.000000     1.000000    -120.000000
25%     1.000000    114.000000    112.000000     1.000000      6.500000
50%     1.000000    162.000000    162.000000     1.000000      9.500000
75%     1.000000    233.000000    233.000000     2.000000     14.500000
max      99.000000    265.000000    265.000000     4.000000     999.990000

```

	extra	mta_tax	tip_amount	tolls_amount \
count	22699.000000	22699.000000	22699.000000	22699.000000
mean	0.333275	0.497445	1.835781	0.312542
std	0.463097	0.039465	2.800626	1.399212
min	-1.000000	-0.500000	0.000000	0.000000
25%	0.000000	0.500000	0.000000	0.000000
50%	0.000000	0.500000	1.350000	0.000000
75%	0.500000	0.500000	2.450000	0.000000
max	4.500000	0.500000	200.000000	19.100000

	improvement_surcharge	total_amount
count	22699.000000	22699.000000
mean	0.299551	16.310502
std	0.015673	16.097295
min	-0.300000	-120.300000
25%	0.300000	8.750000
50%	0.300000	11.800000
75%	0.300000	17.800000
max	0.300000	1200.290000

4.2.2 Task 2b. Convert pickup & dropoff columns to datetime

```
[6]: # Check the format of the data
    ### YOUR CODE HERE ###
    df.dtypes
```

```
[6]: Unnamed: 0          int64
    VendorID           int64
    tpep_pickup_datetime object
    tpep_dropoff_datetime object
    passenger_count     int64
    trip_distance       float64
    RatecodeID         int64
    store_and_fwd_flag  object
    PULocationID       int64
    DOLocationID       int64
    payment_type       int64
    fare_amount        float64
    extra              float64
    mta_tax            float64
    tip_amount         float64
    tolls_amount       float64
    improvement_surcharge float64
    total_amount       float64
    dtype: object
```

```
[7]: # Convert datetime columns to datetime
    ### YOUR CODE HERE ###
    df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
    df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])
```

4.2.3 Task 2c. Create duration column

Create a new column called `duration` that represents the total number of minutes that each taxi ride took.

```
[8]: # Create `duration` column
    ### YOUR CODE HERE ###
    df['duration'] = (df['tpep_dropoff_datetime'] - df['tpep_pickup_datetime'])/np.
        ↳timedelta64(1, 'm')
```

4.2.4 Outliers

Call `df.info()` to inspect the columns and decide which ones to check for outliers.

```
[9]: ### YOUR CODE HERE ###
    df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            22699 non-null  int64
1   VendorID                              22699 non-null  int64
2   tpep_pickup_datetime                  22699 non-null  datetime64[ns]
3   tpep_dropoff_datetime                 22699 non-null  datetime64[ns]
4   passenger_count                       22699 non-null  int64
5   trip_distance                         22699 non-null  float64
6   RatecodeID                           22699 non-null  int64
7   store_and_fwd_flag                    22699 non-null  object
8   PULocationID                         22699 non-null  int64
9   DOLocationID                         22699 non-null  int64
10  payment_type                          22699 non-null  int64
11  fare_amount                           22699 non-null  float64
12  extra                                 22699 non-null  float64
13  mta_tax                               22699 non-null  float64
14  tip_amount                            22699 non-null  float64
15  tolls_amount                          22699 non-null  float64
16  improvement_surcharge                 22699 non-null  float64
17  total_amount                          22699 non-null  float64
18  duration                              22699 non-null  float64
```

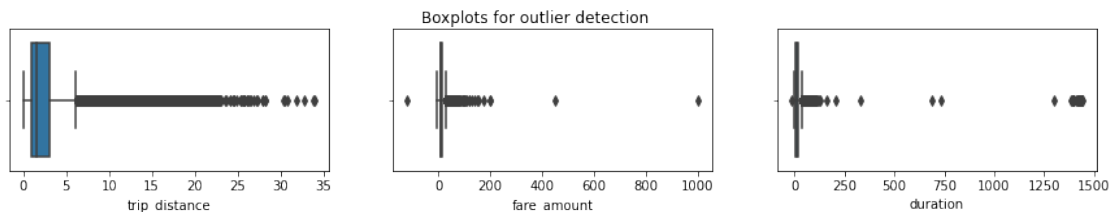
```
dtypes: datetime64[ns](2), float64(9), int64(7), object(1)
memory usage: 3.3+ MB
```

Keeping in mind that many of the features will not be used to fit your model, the most important columns to check for outliers are likely to be: * `trip_distance` * `fare_amount` * `duration`

4.2.5 Task 2d. Box plots

Plot a box plot for each feature: `trip_distance`, `fare_amount`, `duration`.

```
[10]: ### YOUR CODE HERE ###
fig, axes = plt.subplots(1, 3, figsize=(15, 2))
fig.suptitle('Boxplots for outlier detection')
sns.boxplot(ax=axes[0], x=df['trip_distance'])
sns.boxplot(ax=axes[1], x=df['fare_amount'])
sns.boxplot(ax=axes[2], x=df['duration'])
plt.show()
```



Questions: 1. Which variable(s) contains outliers?

2. Are the values in the `trip_distance` column unbelievable?

3. What about the lower end? Do distances, fares, and durations of 0 (or negative values) make sense?

==> ENTER YOUR RESPONSE HERE All three variables contain outliers. Some are extreme, but others not so much.

It's 30 miles from the southern tip of Staten Island to the northern end of Manhattan and that's in a straight line. With this knowledge and the distribution of the values in this column, it's reasonable to leave these values alone and not alter them. However, the values for `fare_amount` and `duration` definitely seem to have problematic outliers on the higher end.

Probably not for the latter two, but for `trip_distance` it might be okay.

4.2.6 Task 2e. Imputations

trip_distance outliers You know from the summary statistics that there are trip distances of 0. Are these reflective of erroneous data, or are they very short trips that get rounded down?

To check, sort the column values, eliminate duplicates, and inspect the least 10 values. Are they rounded values or precise values?

```
[11]: # Are trip distances of 0 bad data or very short trips rounded down?
      ### YOUR CODE HERE ###
      sorted(set(df['trip_distance']))[:10]
```

```
[11]: [0.0, 0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.09]
```

The distances are captured with a high degree of precision. However, it might be possible for trips to have distances of zero if a passenger summoned a taxi and then changed their mind. Besides, are there enough zero values in the data to pose a problem?

Calculate the count of rides where the `trip_distance` is zero.

```
[12]: ### YOUR CODE HERE ###
      sum(df['trip_distance']==0)
```

```
[12]: 148
```

fare_amount outliers

```
[13]: ### YOUR CODE HERE ###
      df['fare_amount'].describe()
```

```
[13]: count      22699.000000
      mean         13.026629
      std          13.243791
      min         -120.000000
      25%           6.500000
      50%           9.500000
      75%          14.500000
      max           999.990000
      Name: fare_amount, dtype: float64
```

Question: What do you notice about the values in the `fare_amount` column? Low values: Negative values are problematic. Values of zero could be legitimate if the taxi logged a trip that was immediately canceled. High values: The maximum fare amount in this dataset is nearly \$1,000, which seems very unlikely. High values for this feature can be capped based on intuition and statistics. The interquartile range (IQR) is \$8. The standard formula of $Q3 + (1.5 * IQR)$ yields \$26.50. That doesn't seem appropriate for the maximum fare cap. In this case, we'll use a factor of 6, which results in a cap of \$62.50.

Impute values less than \$0 with 0.

```
[14]: # Impute values less than $0 with 0
      ### YOUR CODE HERE ###
      df.loc[df['fare_amount'] < 0, 'fare_amount'] = 0
```


Now impute the maximum value as $Q3 + (6 * IQR)$.

```
[19]: def outlier_imputer(column_list, iqr_factor):  
    '''  
    Impute upper-limit values in specified columns based on their interquartile_  
    ↪range.  
  
    Arguments:  
    column_list: A list of columns to iterate over  
    iqr_factor: A number representing x in the formula:  
                 $Q3 + (x * IQR)$ . Used to determine maximum threshold,  
                beyond which a point is considered an outlier.  
  
    The IQR is computed for each column in column_list and values exceeding  
    the upper threshold for each column are imputed with the upper threshold_  
    ↪value.  
    '''  
    for col in column_list:  
        # Reassign minimum to zero  
        df.loc[df[col] < 0, col] = 0  
  
        # Calculate upper threshold  
        q1 = df[col].quantile(0.25)  
        q3 = df[col].quantile(0.75)  
        iqr = q3 - q1  
        upper_threshold = q3 + (iqr_factor * iqr)  
        print(col)  
        print('q3:', q3)  
        print('upper_threshold:', upper_threshold)  
  
        # Reassign values > threshold to threshold  
        df.loc[df[col] > upper_threshold, col] = upper_threshold  
        print(df[col].describe())  
        print()  
    outlier_imputer(['fare_amount'], 6)
```

```
fare_amount  
q3: 14.5  
upper_threshold: 62.5  
count    22699.000000  
mean      12.897913  
std       10.541137  
min        0.000000  
25%        6.500000  
50%        9.500000  
75%       14.500000  
max       62.500000  
Name: fare_amount, dtype: float64
```

duration outliers

```
[20]: # Call .describe() for duration outliers
      ### YOUR CODE HERE ###
      df['duration'].describe()
```

```
[20]: count    22699.000000
      mean      17.013777
      std       61.996482
      min      -16.983333
      25%        6.650000
      50%       11.183333
      75%       18.383333
      max      1439.550000
      Name: duration, dtype: float64
```

The duration column has problematic values at both the lower and upper extremities.

- **Low values:** There should be no values that represent negative time. Impute all negative durations with 0.
- **High values:** Impute high values the same way you imputed the high-end outliers for fares: $Q3 + (6 * IQR)$.

```
[22]: # Impute a 0 for any negative values
      ### YOUR CODE HERE ###
      df.loc[df['duration'] < 0, 'duration'] = 0
```

```
[24]: # Impute the high outliers
      ### YOUR CODE HERE ###
      outlier_imputer(['duration'], 6)
```

```
duration
q3: 18.383333333333333
upper_threshold: 88.78333333333333
count    22699.000000
mean      14.460555
std       11.947043
min        0.000000
25%        6.650000
50%       11.183333
75%       18.383333
max       88.783333
Name: duration, dtype: float64
```

4.2.7 Task 3a. Feature engineering

Create `mean_distance` column When deployed, the model will not know the duration of a trip until after the trip occurs, so you cannot train a model that uses this feature. However, you can use the statistics of trips you *do* know to generalize about ones you do not know.

In this step, create a column called `mean_distance` that captures the mean distance for each group of trips that share pickup and dropoff points.

For example, if your data were:

```
|Trip|Start|End|Distance| |-: |:—:|:~:| | 1 | A | B | 1 | | 2 | C | D | 2 | | 3 | A | B | 1.5 | | 4 | D | C | 3 |
```

The results should be:

```
A -> B: 1.25 miles
C -> D: 2 miles
D -> C: 3 miles
```

Notice that C -> D is not the same as D -> C. All trips that share a unique pair of start and end points get grouped and averaged.

Then, a new column `mean_distance` will be added where the value at each row is the average for all trips with those pickup and dropoff locations:

Trip	Start	End	Distance	mean_distance
1	A	B	1	1.25
2	C	D	2	2
3	A	B	1.5	1.25
4	D	C	3	3

Begin by creating a helper column called `pickup_dropoff`, which contains the unique combination of pickup and dropoff location IDs for each row.

One way to do this is to convert the pickup and dropoff location IDs to strings and join them, separated by a space. The space is to ensure that, for example, a trip with pickup/dropoff points of 12 & 151 gets encoded differently than a trip with points 121 & 51.

So, the new column would look like this:

Trip	Start	End	pickup_dropoff
1	A	B	'A B'
2	C	D	'C D'
3	A	B	'A B'
4	D	C	'D C'

```
[25]: # Create `pickup_dropoff` column
      ### YOUR CODE HERE ###
```

```
df['pickup_dropoff'] = df['PULocationID'].astype(str) + ' ' +
↳df['DOLocationID'].astype(str)
```

Now, use a `groupby()` statement to group each row by the new `pickup_dropoff` column, compute the mean, and capture the values only in the `trip_distance` column. Assign the results to a variable named `grouped`.

```
[30]: ### YOUR CODE HERE ###
grouped = df.groupby(['pickup_dropoff']).mean()[['trip_distance']]
grouped[:5]
```

```
[30]:
```

	trip_distance
pickup_dropoff	
1 1	2.433333
10 148	15.700000
100 1	16.890000
100 100	0.253333
100 107	1.180000

`grouped` is an object of the `DataFrame` class.

1. Convert it to a dictionary using the `to_dict()` method. Assign the results to a variable called `grouped_dict`. This will result in a dictionary with a key of `trip_distance` whose values are another dictionary. The inner dictionary's keys are pickup/dropoff points and its values are mean distances. This is the information you want.

Example:

```
grouped_dict = {'trip_distance': {'A B': 1.25, 'C D': 2, 'D C': 3}}
```

2. Reassign the `grouped_dict` dictionary so it contains only the inner dictionary. In other words, get rid of `trip_distance` as a key, so:

Example:

```
grouped_dict = {'A B': 1.25, 'C D': 2, 'D C': 3}
```

```
[35]: # 1. Convert `grouped` to a dictionary
### YOUR CODE HERE ###
grouped_dict = grouped.to_dict()

# 2. Reassign to only contain the inner dictionary
### YOUR CODE HERE ###
grouped_dict = grouped_dict['trip_distance']
```

1. Create a `mean_distance` column that is a copy of the `pickup_dropoff` helper column.
2. Use the `map()` method on the `mean_distance` series. Pass `grouped_dict` as its argument. Reassign the result back to the `mean_distance` series. When you pass a dictionary to the `Series.map()` method, it will replace the data in the series where that data matches the dictionary's keys. The values that get imputed are the values of the dictionary.

Example:

```
df['mean_distance']
```

mean_distance
'A B'
'C D'
'A B'
'D C'
'E F'

```
grouped_dict = {'A B': 1.25, 'C D': 2, 'D C': 3}
df['mean_distance`'] = df['mean_distance'].map(grouped_dict)
df['mean_distance']
```

mean_distance
1.25
2
1.25
3
NaN

When used this way, the `map()` Series method is very similar to `replace()`, however, note that `map()` will impute NaN for any values in the series that do not have a corresponding key in the mapping dictionary, so be careful.

```
[37]: # 1. Create a mean_distance column that is a copy of the pickup_dropoff helper_
      ↪ column
      ### YOUR CODE HERE ###
      df['mean_distance'] = df['pickup_dropoff']

      # 2. Map `grouped_dict` to the `mean_distance` column
      ### YOUR CODE HERE ###
      df['mean_distance'] = df['mean_distance'].map(grouped_dict)

      # Confirm that it worked
      ### YOUR CODE HERE ###
      df[(df['PULocationID']==100) & (df['DOLocationID']==231)][['mean_distance']]
```

```
[37]:      mean_distance
0      3.521667
4909    3.521667
16636    3.521667
18134    3.521667
19761    3.521667
20581    3.521667
```

Create mean_duration column Repeat the process used to create the mean_distance column to create a mean_duration column.

```
[38]: grouped = df.groupby('pickup_dropoff').mean(numeric_only=True)[['duration']]

# Create a dictionary where keys are unique pickup_dropoffs and values are
# mean trip duration for all trips with those pickup_dropoff combos
grouped_dict = grouped.to_dict()
grouped_dict = grouped_dict['duration']

df['mean_duration'] = df['pickup_dropoff']
df['mean_duration'] = df['mean_duration'].map(grouped_dict)

# Confirm that it worked
df[(df['PULocationID']==100) & (df['DOLocationID']==231)][['mean_duration']]
```

```
[38]:      mean_duration
0      22.847222
4909    22.847222
16636   22.847222
18134   22.847222
19761   22.847222
20581   22.847222
```

Create day and month columns Create two new columns, day (name of day) and month (name of month) by extracting the relevant information from the tpep_pickup_datetime column.

```
[64]: # Create 'day' col
### YOUR CODE HERE ###
df['day'] = df['tpep_pickup_datetime'].dt.day_name().str.lower()

# Create 'month' col
### YOUR CODE HERE ###
df['month'] = df['tpep_pickup_datetime'].dt.month_name().str.lower()
```

Create rush_hour column Define rush hour as: * Any weekday (not Saturday or Sunday) AND * Either from 06:00–10:00 or from 16:00–20:00

Create a binary rush_hour column that contains a 1 if the ride was during rush hour and a 0 if it was not.

```
[65]: # Create 'rush_hour' col
df['rush_hour'] = df['tpep_pickup_datetime'].dt.hour

# If day is Saturday or Sunday, impute 0 in `rush_hour` column
df.loc[df['day'].isin(['saturday', 'sunday']), 'rush_hour'] = 0
```

```
[66]: def rush_hourizer(hour):
        if 6 <= hour['rush_hour'] < 10:
            val = 1
        elif 16 <= hour['rush_hour'] < 20:
            val = 1
        else:
            val = 0
        return val
```

```
[67]: # Apply the `rush_hourizer()` function to the new column
df.loc[(df.day != 'saturday') & (df.day != 'sunday'), 'rush_hour'] = df.
    ↪apply(rush_hourizer, axis=1)
df.head()
```

```
[67]: Unnamed: 0  VendorID tpep_pickup_datetime tpep_dropoff_datetime \
0      24870114          2  2017-03-25 08:55:43  2017-03-25 09:09:47
1      35634249          1  2017-04-11 14:53:28  2017-04-11 15:19:58
2      106203690         1  2017-12-15 07:26:56  2017-12-15 07:34:08
3      38942136          2  2017-05-07 13:17:59  2017-05-07 13:48:14
4      30841670          2  2017-04-15 23:32:20  2017-04-15 23:49:03

    passenger_count  trip_distance  RatecodeID  store_and_fwd_flag \
0                  6           3.34           1                  N
1                  1           1.80           1                  N
2                  1           1.00           1                  N
3                  1           3.70           1                  N
4                  1           4.37           1                  N

    PULocationID  DOLocationID  ...  tolls_amount  improvement_surcharge \
0              100           231  ...           0.0              0.3
1              186           43  ...           0.0              0.3
2              262           236  ...           0.0              0.3
3              188           97  ...           0.0              0.3
4               4           112  ...           0.0              0.3

    total_amount  duration  pickup_dropoff  mean_distance  mean_duration \
0          16.56  14.066667          100 231          3.521667          22.847222
1          20.80  26.500000          186 43          3.108889          24.470370
2           8.75   7.200000          262 236          0.881429           7.250000
3          27.69  30.250000          188 97          3.700000          30.250000
4          17.80  16.716667           4 112          4.435000          14.616667

    day      month  rush_hour
0  saturday   march         0
1   tuesday   april         0
2   friday   december        1
3   sunday     may          0
```

4 saturday april 0

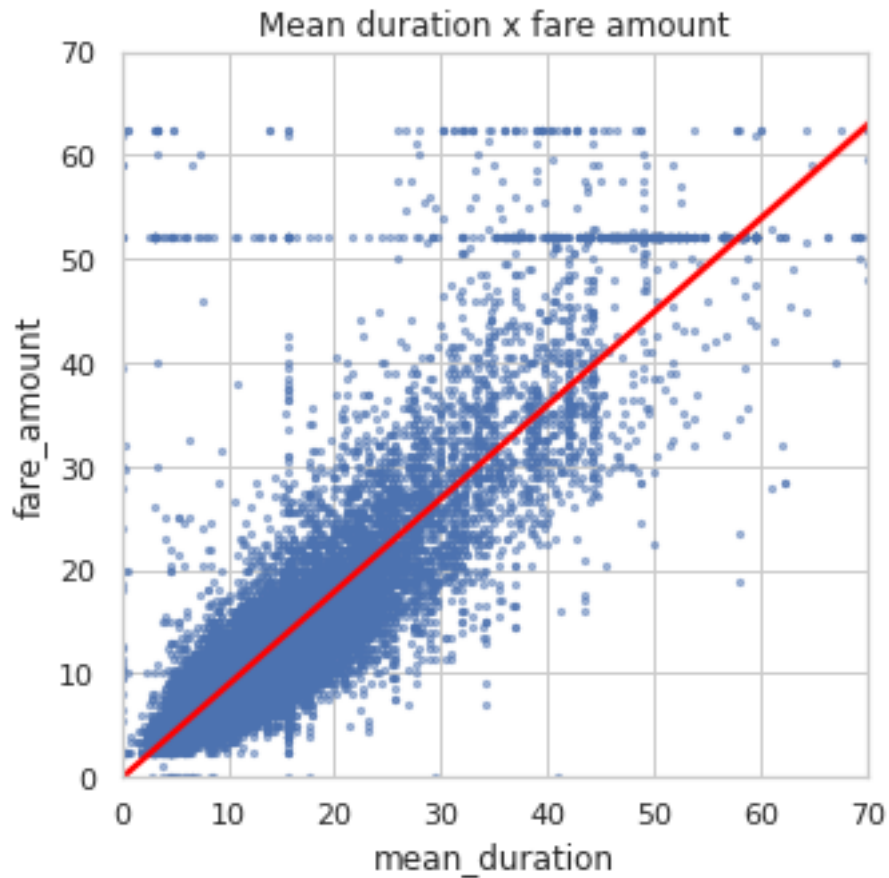
[5 rows x 25 columns]

4.2.8 Task 4. Scatter plot

Create a scatterplot to visualize the relationship between `mean_duration` and `fare_amount`.

```
[69]: # Create a scatterplot to visualize the relationship between variables of interest
      ## YOUR CODE HERE ##
      #sns.scatterplot(x = 'mean_duration', y = 'fare_amount', data = df)

      ###Exempler method
      sns.set(style='whitegrid')
      f = plt.figure()
      f.set_figwidth(5)
      f.set_figheight(5)
      sns.regplot(x=df['mean_duration'], y=df['fare_amount'],
                  scatter_kws={'alpha':0.5, 's':5},
                  line_kws={'color':'red'})
      plt.ylim(0, 70)
      plt.xlim(0, 70)
      plt.title('Mean duration x fare amount')
      plt.show()
```

The `mean_duration` variable correlates with the target variable. But what are the horizontal lines around fare amounts of 52 dollars and 63 dollars? What are the values and how many are there?

You know what one of the lines represents. 62 dollars and 50 cents is the maximum that was imputed for outliers, so all former outliers will now have fare amounts of \$62.50. What is the other line?

Check the value of the rides in the second horizontal line in the scatter plot.

```
[71]: ### YOUR CODE HERE ###
df[df['fare_amount'] > 50]['fare_amount'].value_counts().head()
```

```
[71]: 52.0      514
      62.5       84
      59.0        9
      50.5         9
      57.5         8
      Name: fare_amount, dtype: int64
```

Examine the first 30 of these trips.

```
[72]: # Set pandas to display all columns
      ### YOUR CODE HERE ###
      pd.set_option('display.max_columns', None)
      df[df['fare_amount']==52].head(30)
```

```
[72]:      Unnamed: 0  VendorID tpep_pickup_datetime tpep_dropoff_datetime \
11      18600059          2  2017-03-05 19:15:30  2017-03-05 19:52:18
110     47959795          1  2017-06-03 14:24:57  2017-06-03 15:31:48
161     95729204          2  2017-11-11 20:16:16  2017-11-11 20:17:14
247    103404868          2  2017-12-06 23:37:08  2017-12-07 00:06:19
379     80479432          2  2017-09-24 23:45:45  2017-09-25 00:15:14
388     16226157          1  2017-02-28 18:30:05  2017-02-28 19:09:55
406     55253442          2  2017-06-05 12:51:58  2017-06-05 13:07:35
449     65900029          2  2017-08-03 22:47:14  2017-08-03 23:32:41
468     80904240          2  2017-09-26 13:48:26  2017-09-26 14:31:17
520     33706214          2  2017-04-23 21:34:48  2017-04-23 22:46:23
569     99259872          2  2017-11-22 21:31:32  2017-11-22 22:00:25
572     61050418          2  2017-07-18 13:29:06  2017-07-18 13:29:19
586     54444647          2  2017-06-26 13:39:12  2017-06-26 14:34:54
692     94424289          2  2017-11-07 22:15:00  2017-11-07 22:45:32
717    103094220          1  2017-12-06 05:19:50  2017-12-06 05:53:52
719     66115834          1  2017-08-04 17:53:34  2017-08-04 18:50:56
782     55934137          2  2017-06-09 09:31:25  2017-06-09 10:24:10
816     13731926          2  2017-02-21 06:11:03  2017-02-21 06:59:39
818     52277743          2  2017-06-20 08:15:18  2017-06-20 10:24:37
835      2684305          2  2017-01-10 22:29:47  2017-01-10 23:06:46
840     90860814          2  2017-10-27 21:50:00  2017-10-27 22:35:04
861     106575186         1  2017-12-16 06:39:59  2017-12-16 07:07:59
881     110495611         2  2017-12-30 05:25:29  2017-12-30 06:01:29
958     87017503          1  2017-10-15 22:39:12  2017-10-15 23:14:22
970     12762608          2  2017-02-17 20:39:42  2017-02-17 21:13:29
984     71264442          1  2017-08-23 18:23:26  2017-08-23 19:18:29
1082    11006300          2  2017-02-07 17:20:19  2017-02-07 17:34:41
1097    68882036          2  2017-08-14 23:01:15  2017-08-14 23:03:35
1110    74720333          1  2017-09-06 10:46:17  2017-09-06 11:44:41
1179    51937907          2  2017-06-19 06:23:13  2017-06-19 07:03:53
```

```
      passenger_count  trip_distance  RatecodeID  store_and_fwd_flag \
11                  2          18.90          2                N
110                 1          18.00          2                N
161                 1           0.23          2                N
247                 1          18.93          2                N
379                 1          17.99          2                N
388                 1          18.40          2                N
406                 1           4.73          2                N
449                 2          18.21          2                N
468                 1          17.27          2                N
```

520	6	18.34	2	N
569	1	18.65	2	N
572	1	0.00	2	N
586	1	17.76	2	N
692	2	16.97	2	N
717	1	20.80	2	N
719	1	21.60	2	N
782	2	18.81	2	N
816	5	16.94	2	N
818	1	17.77	2	N
835	1	18.57	2	N
840	1	22.43	2	N
861	2	17.80	2	N
881	6	18.23	2	N
958	1	21.80	2	N
970	1	19.57	2	N
984	1	16.70	2	N
1082	1	1.09	2	N
1097	5	2.12	2	N
1110	1	19.10	2	N
1179	6	19.77	2	N

	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	\
11	236	132	1	52.0	0.0	0.5	
110	132	163	1	52.0	0.0	0.5	
161	132	132	2	52.0	0.0	0.5	
247	132	79	2	52.0	0.0	0.5	
379	132	234	1	52.0	0.0	0.5	
388	132	48	2	52.0	4.5	0.5	
406	228	88	2	52.0	0.0	0.5	
449	132	48	2	52.0	0.0	0.5	
468	186	132	2	52.0	0.0	0.5	
520	132	148	1	52.0	0.0	0.5	
569	132	144	1	52.0	0.0	0.5	
572	230	161	1	52.0	0.0	0.5	
586	211	132	1	52.0	0.0	0.5	
692	132	170	1	52.0	0.0	0.5	
717	132	239	1	52.0	0.0	0.5	
719	264	264	1	52.0	4.5	0.5	
782	163	132	1	52.0	0.0	0.5	
816	132	170	1	52.0	0.0	0.5	
818	132	246	1	52.0	0.0	0.5	
835	132	48	1	52.0	0.0	0.5	
840	132	163	2	52.0	0.0	0.5	
861	75	132	1	52.0	0.0	0.5	
881	68	132	2	52.0	0.0	0.5	
958	132	261	2	52.0	0.0	0.5	

970	132	140	1	52.0	0.0	0.5
984	132	230	1	52.0	4.5	0.5
1082	170	48	2	52.0	4.5	0.5
1097	265	265	2	52.0	0.0	0.5
1110	239	132	1	52.0	0.0	0.5
1179	238	132	1	52.0	0.0	0.5

	tip_amount	tolls_amount	improvement_surcharge	total_amount	\
11	14.58	5.54	0.3	72.92	
110	0.00	0.00	0.3	52.80	
161	0.00	0.00	0.3	52.80	
247	0.00	0.00	0.3	52.80	
379	14.64	5.76	0.3	73.20	
388	0.00	5.54	0.3	62.84	
406	0.00	5.76	0.3	58.56	
449	0.00	5.76	0.3	58.56	
468	0.00	5.76	0.3	58.56	
520	5.00	0.00	0.3	57.80	
569	10.56	0.00	0.3	63.36	
572	11.71	5.76	0.3	70.27	
586	11.71	5.76	0.3	70.27	
692	11.71	5.76	0.3	70.27	
717	5.85	5.76	0.3	64.41	
719	12.60	5.76	0.3	75.66	
782	13.20	0.00	0.3	66.00	
816	2.00	5.54	0.3	60.34	
818	11.71	5.76	0.3	70.27	
835	13.20	0.00	0.3	66.00	
840	0.00	5.76	0.3	58.56	
861	6.00	5.76	0.3	64.56	
881	0.00	0.00	0.3	52.80	
958	0.00	0.00	0.3	52.80	
970	11.67	5.54	0.3	70.01	
984	42.29	0.00	0.3	99.59	
1082	0.00	5.54	0.3	62.84	
1097	0.00	0.00	0.3	52.80	
1110	15.80	0.00	0.3	68.60	
1179	17.57	5.76	0.3	76.13	

	duration	pickup_dropoff	mean_distance	mean_duration	day	\
11	36.800000	236 132	19.211667	40.500000	sunday	
110	66.850000	132 163	19.229000	52.941667	saturday	
161	0.966667	132 132	2.255862	3.021839	saturday	
247	29.183333	132 79	19.431667	47.275000	wednesday	
379	29.483333	132 234	17.654000	49.833333	sunday	
388	39.833333	132 48	18.761905	58.246032	tuesday	
406	15.616667	228 88	4.730000	15.616667	monday	

449	45.450000	132 48	18.761905	58.246032	thursday
468	42.850000	186 132	17.096000	42.920000	tuesday
520	71.583333	132 148	17.994286	46.340476	sunday
569	28.883333	132 144	18.537500	37.000000	wednesday
572	0.216667	230 161	0.685484	7.965591	tuesday
586	55.700000	211 132	16.580000	61.691667	monday
692	30.533333	132 170	17.203000	37.113333	tuesday
717	34.033333	132 239	20.901250	44.862500	wednesday
719	57.366667	264 264	3.191516	15.618773	friday
782	52.750000	163 132	17.275833	52.338889	friday
816	48.600000	132 170	17.203000	37.113333	tuesday
818	88.783333	132 246	18.515000	66.316667	tuesday
835	36.983333	132 48	18.761905	58.246032	tuesday
840	45.066667	132 163	19.229000	52.941667	friday
861	28.000000	75 132	18.442500	36.204167	saturday
881	36.000000	68 132	18.785000	58.041667	saturday
958	35.166667	132 261	22.115000	51.493750	sunday
970	33.783333	132 140	19.293333	36.791667	friday
984	55.050000	132 230	18.571200	59.598000	wednesday
1082	14.366667	170 48	1.265789	14.135965	tuesday
1097	2.333333	265 265	0.753077	3.411538	monday
1110	58.400000	239 132	19.795000	50.562500	wednesday
1179	40.666667	238 132	19.470000	53.861111	monday

	month	rush_hour
11	march	0
110	june	0
161	november	0
247	december	0
379	september	0
388	february	1
406	june	0
449	august	0
468	september	0
520	april	0
569	november	0
572	july	0
586	june	0
692	november	0
717	december	0
719	august	1
782	june	1
816	february	1
818	june	1
835	january	0
840	october	0
861	december	0

881	december	0
958	october	0
970	february	0
984	august	1
1082	february	1
1097	august	0
1110	september	0
1179	june	1

Question: What do you notice about the first 30 trips?

==> ENTER YOUR RESPONSE HERE It seems that almost all of the trips in the first 30 rows where the fare amount was \$52 either begin or end at location 132, and all of them have a RatecodeID of 2.

There is no readily apparent reason why PULocation 132 should have so many fares of 52 dollars. They seem to occur on all different days, at different times, with both vendors, in all months. However, there are many toll amounts of \$5.76 and \$5.54. This would seem to indicate that location 132 is in an area that frequently requires tolls to get to and from. It's likely this is an airport.

The data dictionary says that RatecodeID of 2 indicates trips for JFK, which is John F. Kennedy International Airport. A quick Google search for “new york city taxi flat rate \$52” indicates that in 2017 (the year that this data was collected) there was indeed a flat fare for taxi trips between JFK airport (in Queens) and Manhattan.

Because RatecodeID is known from the data dictionary, the values for this rate code can be imputed back into the data after the model makes its predictions. This way you know that those data points will always be correct.

4.2.9 Task 5. Isolate modeling variables

Drop features that are redundant, irrelevant, or that will not be available in a deployed environment.

```
[73]: ### YOUR CODE HERE ###
df2 = df.copy()

df2 = df2.drop(['Unnamed: 0', 'tpep_dropoff_datetime', 'tpep_pickup_datetime',
               'trip_distance', 'RatecodeID', 'store_and_fwd_flag',
               ↪ 'PULocationID', 'DOLocationID',
               'payment_type', 'extra', 'mta_tax', 'tip_amount',
               ↪ 'tolls_amount', 'improvement_surcharge',
               'total_amount', 'tpep_dropoff_datetime', 'tpep_pickup_datetime',
               ↪ 'duration',
               'pickup_dropoff', 'day', 'month'
               ], axis=1)

df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   VendorID              22699 non-null  int64
1   passenger_count       22699 non-null  int64
2   fare_amount           22699 non-null  float64
3   mean_distance         22699 non-null  float64
4   mean_duration         22699 non-null  float64
5   rush_hour             22699 non-null  int64
dtypes: float64(3), int64(3)
memory usage: 1.0 MB

```

```
[ ]: ### YOUR CODE HERE ###
```

4.2.10 Task 6. Pair plot

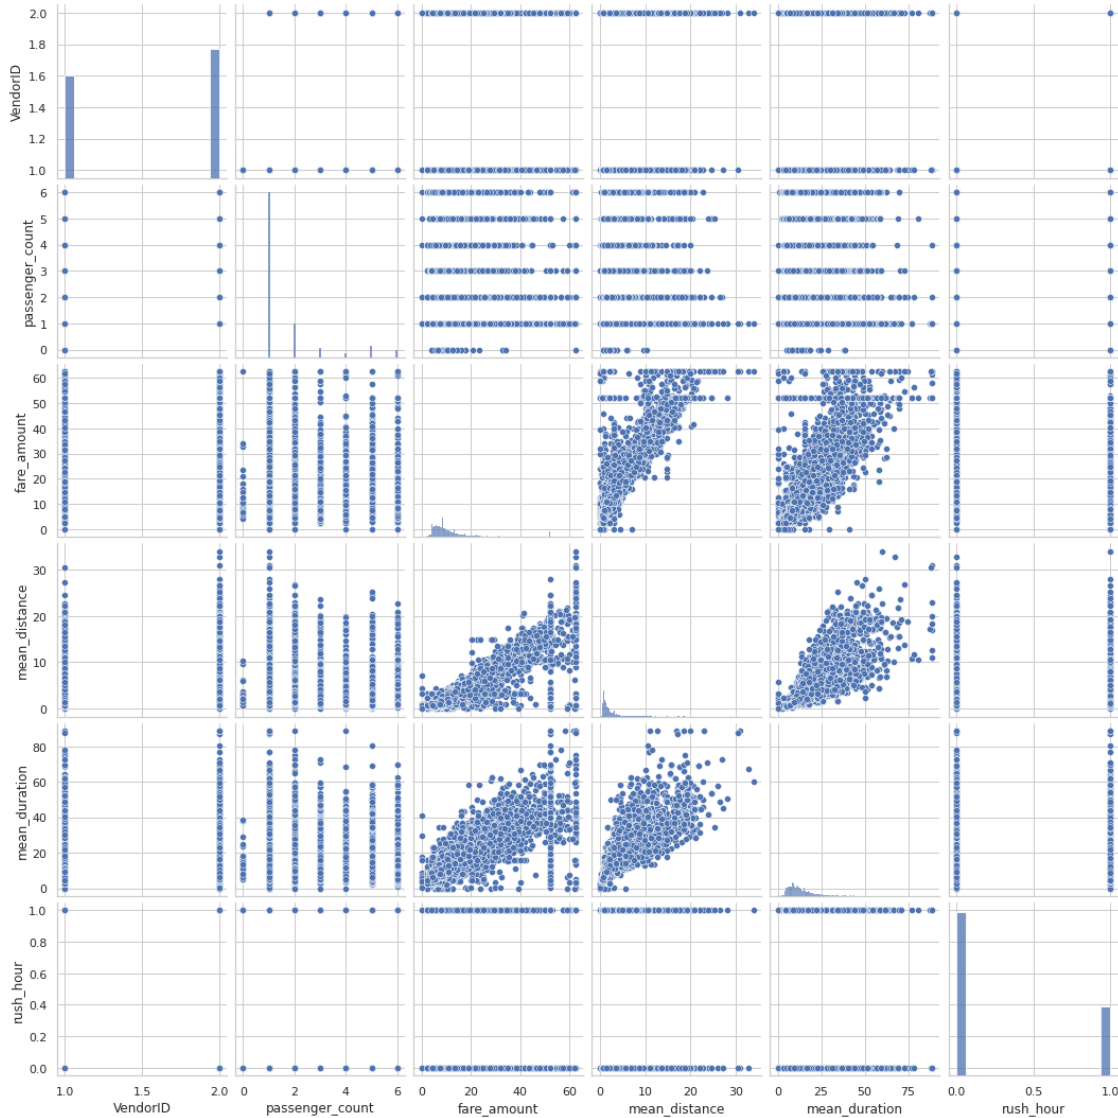
Create a pairplot to visualize pairwise relationships between `fare_amount`, `mean_duration`, and `mean_distance`.

```

[75]: # Create a pairplot to visualize pairwise relationships between variables in
      ↳ the data
      ### YOUR CODE HERE ###
      sns.pairplot(df2)

```

```
[75]: <seaborn.axisgrid.PairGrid at 0x7f26d6c74f10>
```



These variables all show linear correlation with each other. Investigate this further.

4.2.11 Task 7. Identify correlations

Next, code a correlation matrix to help determine most correlated variables.

```
[76]: # Correlation matrix to help determine most correlated variables
      ### YOUR CODE HERE ###
      df2.corr(method='pearson')
```

```
[76]:
```

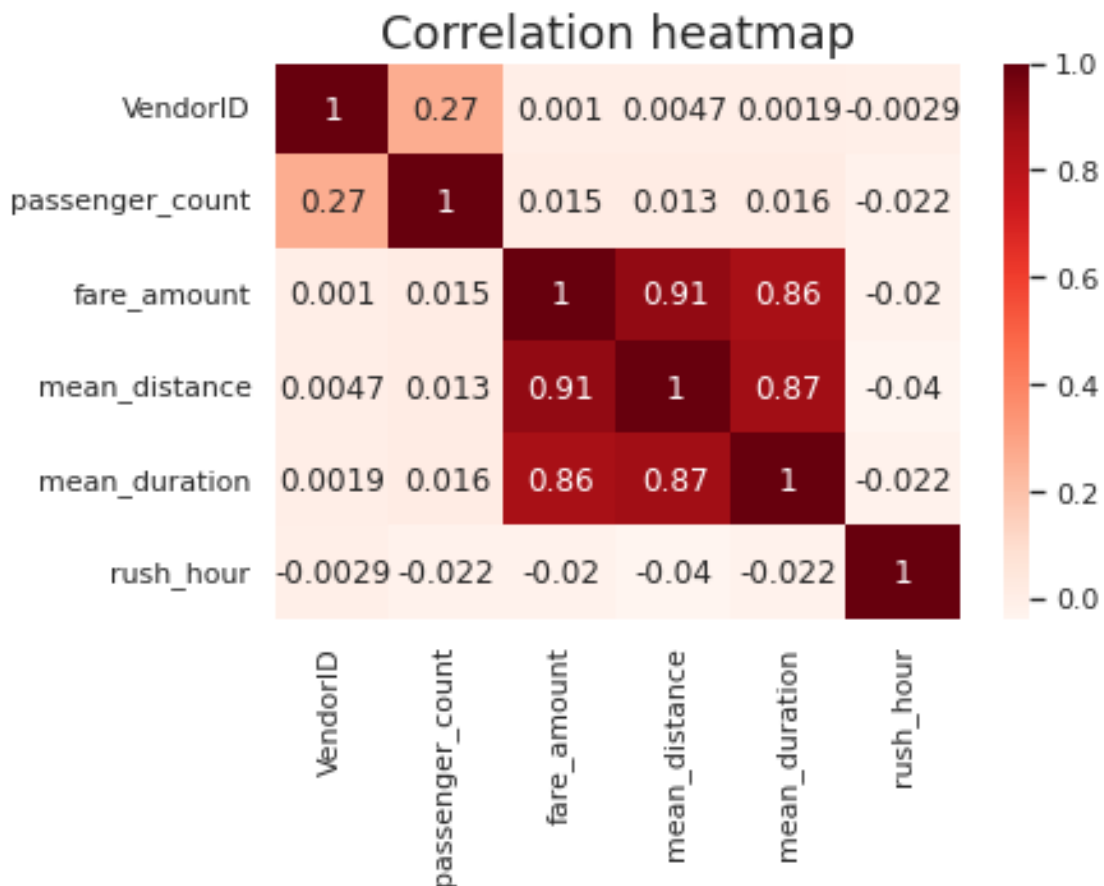
	VendorID	passenger_count	fare_amount	mean_distance	\
VendorID	1.000000	0.266463	0.001045	0.004741	
passenger_count	0.266463	1.000000	0.014942	0.013428	

fare_amount	0.001045	0.014942	1.000000	0.910185
mean_distance	0.004741	0.013428	0.910185	1.000000
mean_duration	0.001876	0.015852	0.859105	0.874864
rush_hour	-0.002874	-0.022035	-0.020075	-0.039725

	mean_duration	rush_hour
VendorID	0.001876	-0.002874
passenger_count	0.015852	-0.022035
fare_amount	0.859105	-0.020075
mean_distance	0.874864	-0.039725
mean_duration	1.000000	-0.021583
rush_hour	-0.021583	1.000000

Visualize a correlation heatmap of the data.

```
[77]: # Create correlation heatmap
      ### YOUR CODE HERE ###
      plt.figure(figsize=(6,4))
      sns.heatmap(df2.corr(method='pearson'), annot=True, cmap='Reds')
      plt.title('Correlation heatmap',
                fontsize=18)
      plt.show()
```



Question: Which variable(s) are correlated with the target variable of `fare_amount`? `mean_duration` and `mean_distance` are both highly correlated with the target variable of `fare_amount`. They're also both correlated with each other, with a Pearson correlation of 0.87.

Recall that highly correlated predictor variables can be bad for linear regression models when you want to be able to draw statistical inferences about the data from the model. However, correlated predictor variables can still be used to create an accurate predictor if the prediction itself is more important than using the model as a tool to learn about your data.

Try modeling with both variables even though they are correlated.

4.3 PACE: Construct

After analysis and deriving variables with close relationships, it is time to begin constructing the model. Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

4.3.1 Task 8a. Split data into outcome variable and features

```
[78]: ### YOUR CODE HERE ###
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 6 columns):
 #   Column            Non-Null Count  Dtype
---  -
 0   VendorID          22699 non-null  int64
 1   passenger_count    22699 non-null  int64
 2   fare_amount        22699 non-null  float64
 3   mean_distance      22699 non-null  float64
 4   mean_duration      22699 non-null  float64
 5   rush_hour          22699 non-null  int64
dtypes: float64(3), int64(3)
memory usage: 1.0 MB
```

Set your X and y variables. X represents the features and y represents the outcome (target) variable.

```
[79]: # Remove the target column from the features
X = df2.drop(columns=['fare_amount'])

# Set y variable
y = df2[['fare_amount']]

# Display first few rows
X.head()
```

```
[79]:
```

	VendorID	passenger_count	mean_distance	mean_duration	rush_hour
0	2	6	3.521667	22.847222	0
1	1	1	3.108889	24.470370	0
2	1	1	0.881429	7.250000	1
3	2	1	3.700000	30.250000	0
4	2	1	4.435000	14.616667	0

4.3.2 Task 8b. Pre-process data

Dummy encode categorical variables

```
[80]: # Convert VendorID to string
      ### YOUR CODE HERE ###
      X['VendorID'] = X['VendorID'].astype(str)

      # Get dummies
      ### YOUR CODE HERE ###
      X = pd.get_dummies(X, drop_first=True)
      X.head()
```

```
[80]:
```

	passenger_count	mean_distance	mean_duration	rush_hour	VendorID_2
0	6	3.521667	22.847222	0	1
1	1	3.108889	24.470370	0	0
2	1	0.881429	7.250000	1	0
3	1	3.700000	30.250000	0	1
4	1	4.435000	14.616667	0	1

4.3.3 Split data into training and test sets

Create training and testing sets. The test set should contain 20% of the total samples. Set random_state=0.

```
[81]: # Create training and testing sets
      ##### YOUR CODE HERE #####
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=0)
```

4.3.4 Standardize the data

Use StandardScaler(), fit(), and transform() to standardize the X_train variables. Assign the results to a variable called X_train_scaled.

```
[82]: # Standardize the X variables
      ### YOUR CODE HERE ###
      scaler = StandardScaler().fit(X_train)
```

```
X_train_scaled = scaler.transform(X_train)
print('X_train scaled:', X_train_scaled)
```

```
X_train scaled: [[-0.50301524  0.8694684   0.17616665 -0.64893329  0.89286563]
 [-0.50301524 -0.60011281 -0.69829589  1.54099045  0.89286563]
 [ 0.27331093 -0.47829156 -0.57301906 -0.64893329 -1.11998936]
 ...
 [-0.50301524 -0.45121122 -0.6788917   -0.64893329 -1.11998936]
 [-0.50301524 -0.58944763 -0.85743597  1.54099045 -1.11998936]
 [ 1.82596329  0.83673851  1.13212101 -0.64893329  0.89286563]]
```

4.3.5 Fit the model

Instantiate your model and fit it to the training data.

```
[83]: # Fit your model to the training data
      ### YOUR CODE HERE ###
      lr=LinearRegression()
      lr.fit(X_train_scaled, y_train)
```

```
[83]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

4.3.6 Task 8c. Evaluate model

4.3.7 Train data

Evaluate your model performance by calculating the residual sum of squares and the explained variance score (R^2). Calculate the Mean Absolute Error, Mean Squared Error, and the Root Mean Squared Error.

```
[86]: # Evaluate the model performance on the training data
      ### YOUR CODE HERE ###
      r_sq = lr.score(X_train_scaled, y_train)
      print('Coefficient of determination:', r_sq)
      y_pred_train = lr.predict(X_train_scaled)
      print('R^2:', r2_score(y_train, y_pred_train))
      print('MAE:', mean_absolute_error(y_train, y_pred_train))
      print('MSE:', mean_squared_error(y_train, y_pred_train))
      print('RMSE:', np.sqrt(mean_squared_error(y_train, y_pred_train)))
```

```
Coefficient of determination: 0.8398434585044773
R^2: 0.8398434585044773
MAE: 2.186666416775414
MSE: 17.88973296349268
RMSE: 4.229625629236313
```

4.3.8 Test data

Calculate the same metrics on the test data. Remember to scale the `X_test` data using the scaler that was fit to the training data. Do not refit the scaler to the testing data, just transform it. Call the results `X_test_scaled`.

```
[84]: # Scale the X_test data
      ### YOUR CODE HERE ###
      X_test_scaled = scaler.transform(X_test)

[85]: # Evaluate the model performance on the testing data
      ### YOUR CODE HERE ###
      r_sq_test = lr.score(X_test_scaled, y_test)
      print('Coefficient of determination:', r_sq_test)
      y_pred_test = lr.predict(X_test_scaled)
      print('R^2:', r2_score(y_test, y_pred_test))
      print('MAE:', mean_absolute_error(y_test, y_pred_test))
      print('MSE:', mean_squared_error(y_test, y_pred_test))
      print('RMSE:', np.sqrt(mean_squared_error(y_test, y_pred_test)))
```

```
Coefficient of determination: 0.8682583641795454
R^2: 0.8682583641795454
MAE: 2.1336549840593864
MSE: 14.326454156998944
RMSE: 3.785030271609323
```

4.4 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

4.4.1 Task 9a. Results

Use the code cell below to get `actual`, `predicted`, and `residual` for the testing set, and store them as columns in a `results` dataframe.

```
[87]: # Create a `results` dataframe
      ### YOUR CODE HERE ###
      results = pd.DataFrame(data={'actual': y_test['fare_amount'],
                                   'predicted': y_pred_test.ravel()})
      results['residual'] = results['actual'] - results['predicted']
      results.head()
```

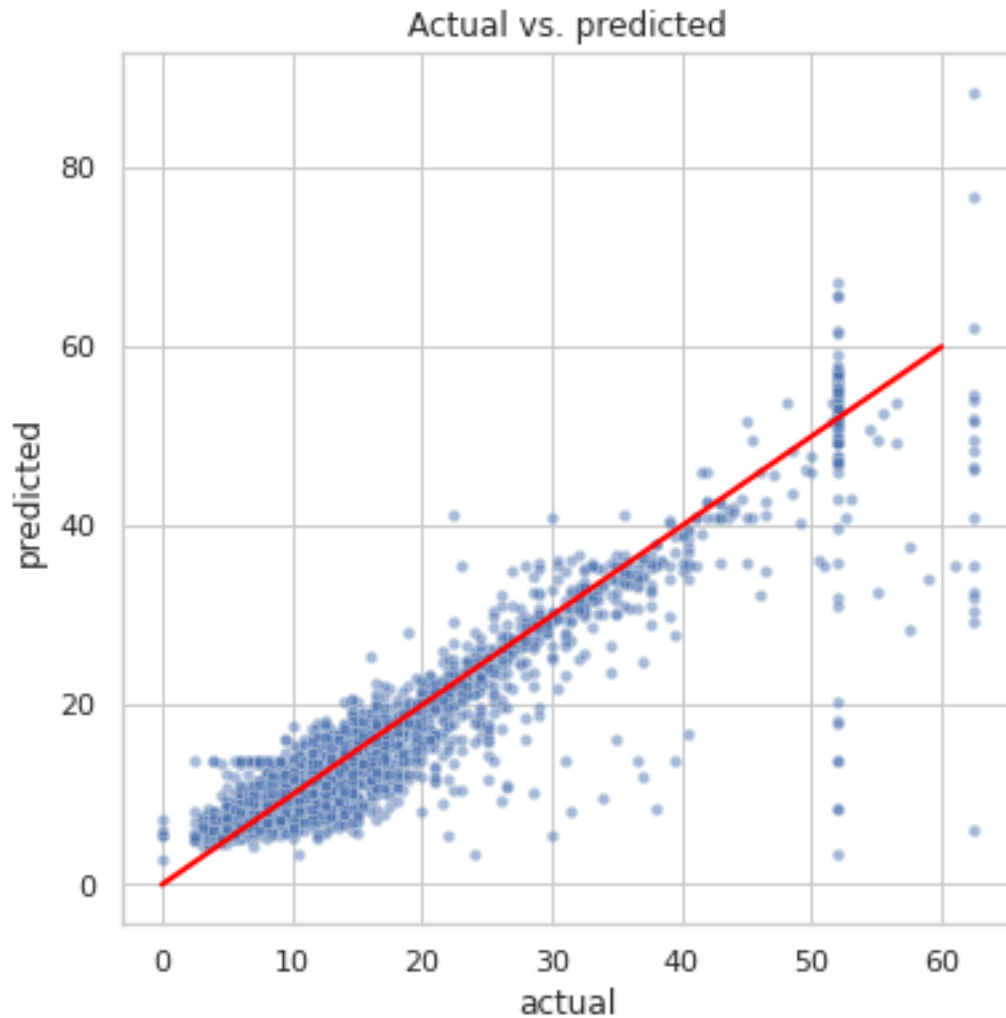
```
[87]:      actual  predicted  residual
5818     14.0    12.356503    1.643497
18134    28.0    16.314595   11.685405
4655      5.5     6.726789   -1.226789
7378     15.5    16.227206   -0.727206
```

13914 9.5 10.536408 -1.036408

4.4.2 Task 9b. Visualize model results

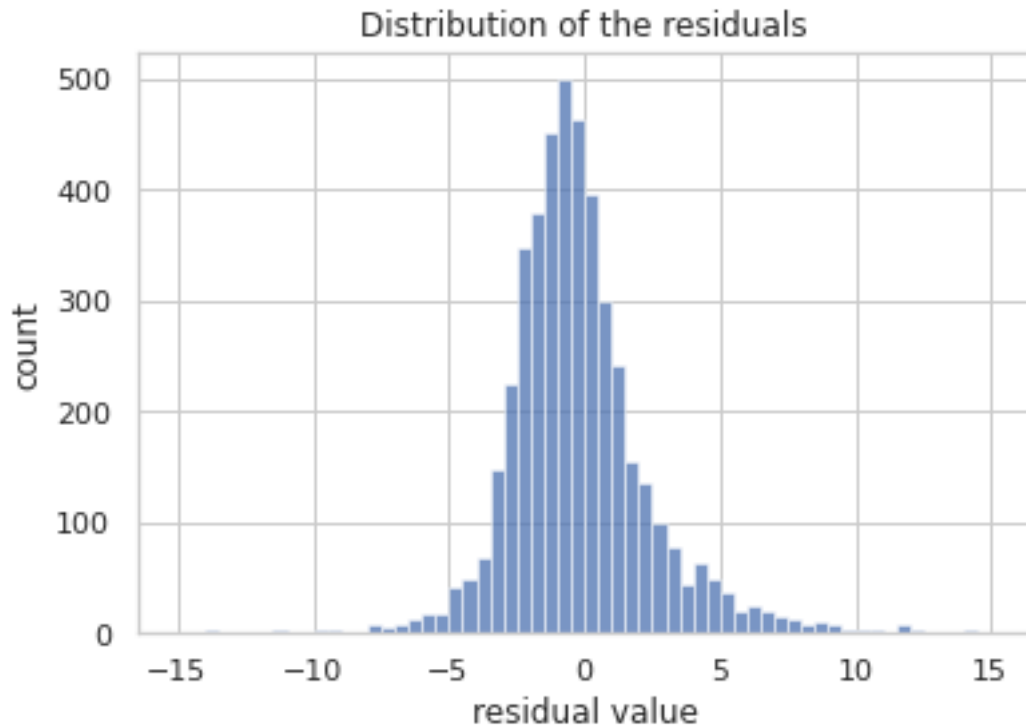
Create a scatterplot to visualize actual vs. predicted.

```
[88]: # Create a scatterplot to visualize `predicted` over `actual`
      ### YOUR CODE HERE ###
      fig, ax = plt.subplots(figsize=(6, 6))
      sns.set(style='whitegrid')
      sns.scatterplot(x='actual',
                     y='predicted',
                     data=results,
                     s=20,
                     alpha=0.5,
                     ax=ax
      )
      # Draw an x=y line to show what the results would be if the model were perfect
      plt.plot([0,60], [0,60], c='red', linewidth=2)
      plt.title('Actual vs. predicted');
```



Visualize the distribution of the `residuals` using a histogram.

```
[89]: # Visualize the distribution of the `residuals`  
      ### YOUR CODE HERE ###  
      sns.histplot(results['residual'], bins=np.arange(-15,15.5,0.5))  
      plt.title('Distribution of the residuals')  
      plt.xlabel('residual value')  
      plt.ylabel('count');
```

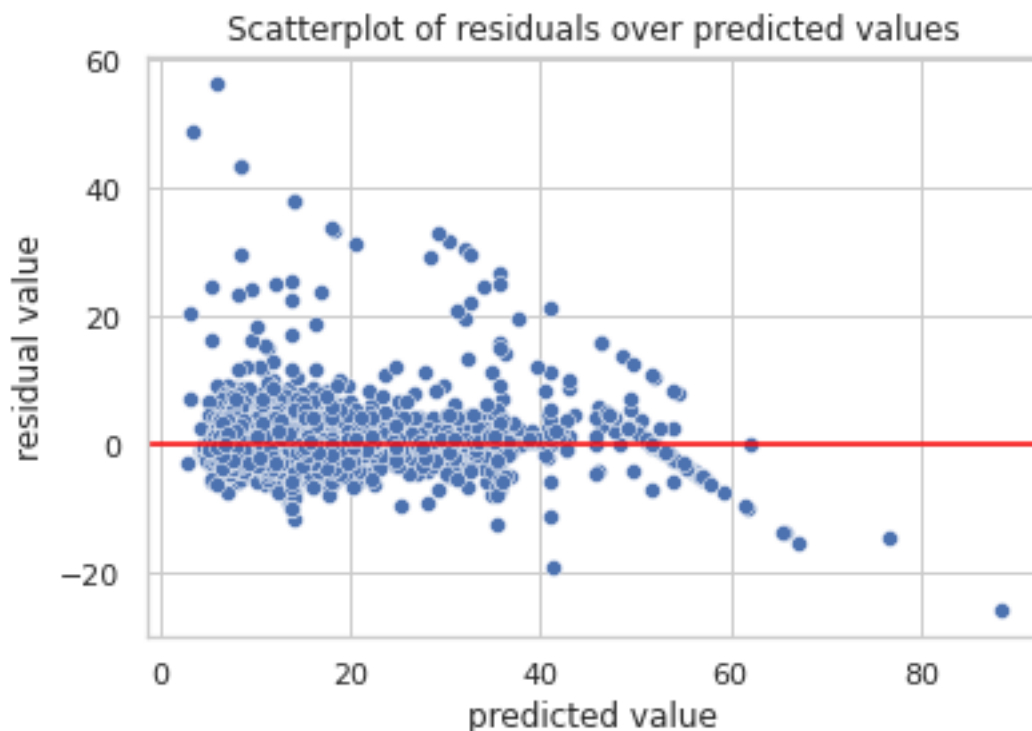


```
[90]: # Calculate residual mean
      ### YOUR CODE HERE ###
      results['residual'].mean()
```

```
[90]: -0.01544262152868053
```

Create a scatterplot of residuals over predicted.

```
[91]: # Create a scatterplot of `residuals` over `predicted`
      ### YOUR CODE HERE ###
      sns.scatterplot(x='predicted', y='residual', data=results)
      plt.axhline(0, c='red')
      plt.title('Scatterplot of residuals over predicted values')
      plt.xlabel('predicted value')
      plt.ylabel('residual value')
      plt.show()
```

4.4.3 Task 9c. Coefficients

Use the `coef_` attribute to get the model's coefficients. The coefficients are output in the order of the features that were used to train the model. Which feature had the greatest effect on trip fare?

```
[92]: # Output the model's coefficients
coefficients = pd.DataFrame(lr.coef_, columns=X.columns)
coefficients
```

```
[92]:
```

	passenger_count	mean_distance	mean_duration	rush_hour	VendorID_2
0	0.030825	7.133867	2.812115	0.110233	-0.054373

What do these coefficients mean? How should they be interpreted?

==> ENTER YOUR RESPONSE HERE The coefficients reveal that `mean_distance` was the feature with the greatest weight in the model's final prediction. Be careful here! A common misinterpretation is that for every mile traveled, the fare amount increases by a mean of \$7.13. This is incorrect. Remember, the data used to train the model was standardized with `StandardScaler()`. As such, the units are no longer miles. In other words, you cannot say "for every mile traveled...", as stated above. The correct interpretation of this coefficient is: controlling for other variables, for every +1 change in standard deviation, the fare amount increases by a mean of \$7.13.

Note also that because some highly correlated features were not removed, the confidence interval of this assessment is wider.

So, translate this back to miles instead of standard deviation (i.e., unscale the data).

Calculate the standard deviation of mean_distance in the X_train data.

Divide the coefficient (7.133867) by the result to yield a more intuitive interpretation.

```
[93]: # 1. Calculate SD of `mean_distance` in X_train data
print(X_train['mean_distance'].std())

# 2. Divide the model coefficient by the standard deviation
print(7.133867 / X_train['mean_distance'].std())
```

```
3.574812975256415
1.9955916713344426
```

```
[ ]: Now you can make a more intuitive interpretation: for every 3.57 miles traveled,
the fare increased by a mean of $7.13. Or, reduced: for every 1 mile traveled,
the fare increased by a mean of $2.00.
```

4.4.4 Task 9d. Conclusion

1. What are the key takeaways from this notebook? Multiple linear regression is a powerful tool to estimate a dependent continuous variable from several independent variables. Exploratory data analysis is useful for selecting both numeric and categorical features for multiple linear regression. Fitting multiple linear regression models may require trial and error to select variables that fit an accurate model while maintaining model assumptions (or not, depending on your use case).
2. What results can be presented from this notebook? You can discuss meeting linear regression assumptions, and you can present the MAE and RMSE scores obtained from the model.

==> ENTER YOUR RESPONSE HERE

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.